

MEASURED DOCUMENT AUTOMATION

Which documents can skip review?

Gemini extracts the order. Deterministic checks decide whether a person still needs to see it.

Example partner rule Accept only when advertiser, order ID, flight dates, and explicit gross amount are present and valid.

9/12

bypassed header review
All 9 had the correct advertiser, order ID, and gross amount.

3.7s median decision time	100% accepted key fields correct	\$0.0231 total run cost / accepted doc
------------------------------	-------------------------------------	---

WHERE THE LINE SITS


Header fields
Promising enough for a controlled pilot
The gate removes routine review while preserving a visible queue.

Line items
Keep human review
29.1% exact-row F1 across 107 trustworthy rows. One wrong field makes the whole row incorrect.

SYSTEM BOUNDARY

Keep business decisions outside the model

Gemini reads the source. Versioned code normalizes values, checks the contract, routes the document, and measures the result.



OPERATING MEASURE

Measure review saved without hiding accepted errors

How much review disappears at an acceptable error rate, and where do the remaining errors cluster?

PAIRED MODEL RUN				Same 12 documents · same prompt · same gate	
Primary	Auto-accept	Accepted key fields	Total cost / accepted		
Gemini 3.5 Flash	75.0%	100%	\$0.0231	USE	
Cost challenger	Auto-accept	Accepted key fields	Total cost / accepted		
Gemini 3.1 Flash-Lite	66.7%	87.5%	\$0.0042	HOLD	
Flash-Lite is cheaper, but one wrong critical header bypassed review. That breaks the operating rule.					

VRDU AD-BUY FORMS DEMO SET

123107

unseen-template documents

pages across dense, rotated, OCR, and financial-edge slices

trustworthy line-item labels scored

Five conflicting headers and 64 corrupted row labels are excluded and recorded.

BEFORE PRODUCTION

Expand to 100-200 jointly labeled documents

Stratify by failure slice, add confidence intervals, and choose the gate with the partner.

Monitor: golden-set regression, null and review drift, reviewer corrections, and random audits of accepted documents.

Schema valid	Header field pass	Exact row F1	Aligned row fields
100%	88.3%	29.1%	79.2%
		31 exact · 106 extracted / 107 reference rows	301 of 380 fields

OBSERVED DURING THE BUILD

Friction worth fixing

Reproduced problem, partner impact, workaround, and the product change that would remove it.

ISSUE	EVIDENCE AND PARTNER IMPACT	WORKAROUND	PRODUCT ASK
P1 · 30 MIN AI Studio Build RPC	The caller does not have permission on three models after a paid Tier 1 key was selected. The UI did not identify whether the problem was the key, project, model, or builder.	Built the Next.js prototype locally with the same Gemini key; kept AI Studio for prompt and schema work.	Show the denied principal, missing permission, and a direct fix in the error.
P1 · 20 MIN PDF resolution controls	resolution and media_resolution on Interactions PDF items returned 400. Service behavior, SDK types, and document guidance did not agree.	Rendered pages at 180 DPI and sent image items with resolution: high .	Align the API, SDK, and docs. Add one executable high-resolution PDF example.
P1 · QUALITY Valid JSON, wrong value	Flash-Lite returned schema-valid output and still auto-accepted one wrong critical header. Shape validation alone can create silent errors.	Kept normalization, semantic checks, routing, benchmark comparison, and audit outside the model.	Publish a reference extractor that covers structured output, semantic evaluation, and monitoring.
P1 · 15 MIN Field contract and labels	Invoice #, Order #, and Net Total were all plausible until the target contract was written. Source audit also found five conflicting headers and 64 corrupted row labels.	Defined field precedence, recorded adjudications, and excluded disputed labels from headline metrics.	Ship field definitions, annotation provenance, and label-quality flags with extraction benchmarks.

WHAT PARTNERS NEED

A useful reference should show what happens after the model returns JSON.