

MEDIA-BUY OPERATIONS

Which documents can skip review?

Turn messy order PDFs into a structured record. Gemini reads the page; code decides whether a person still needs to look.

Example partner rule Auto-accept headers only when advertiser, order ID, flight dates, and explicit gross amount are present and valid.

9/12

skipped header review

All 9 still matched the reference on advertiser, order ID, and gross.

4.1s

median decision time

100%

accepted headers correct

\$0.0234

workload cost / accepted doc

WHERE THE LINE SITS

Headers

Ready for a controlled pilot

Routine header review drops; incomplete or invalid headers stay in a visible queue.

Spots


Keep human review

59% of reference spots fully matched. One wrong field fails the whole spot. Not ready to auto-post.

SYSTEM BOUNDARY

Keep business decisions outside the model

Gemini reads the source. Versioned code normalizes values, checks the field contract, routes the document, and measures the result.



OPERATING MEASURE

Measure review saved without hiding accepted errors

Three questions: how much review disappears, how often accept is actually right, and which failure modes remain.

PAIRED MODEL RUN				Same 12 documents · same prompt · same gate	
Primary	Review saved	Accepts correct	Cost / accepted	USE	
Gemini 3.5 Flash	75%	100%	\$0.0234		
Cost challenger	Review saved	Accepts correct	Cost / accepted	HOLD	
Gemini 3.1 Flash-Lite	75%	100%	\$0.0038		
Flash-Lite is cheaper and matched this run. Held because a prior run let one wrong critical header through — under this policy, that is a disqualifier until a larger paired run.					
FAILURE MODES WE SCORE					
WRONG FIELD		Invoice # or Net Total chosen instead of Order # / Gross			
OVERFILL		Model invents a value where the reference is null			
SPOT MISS		One bad field fails the whole placement			
5 conflicting headers and 64 corrupted spot labels excluded and recorded.					
BEFORE PRODUCTION					
Expand to 100–200 jointly labeled docs					
Stratify by failure mode, add confidence intervals, and set the gate with the partner.					
Monitor: accept-error rate, review rate drift, null-rate drift, reviewer corrections, random audits of accepted docs.					
Schema valid	Header fields correct	Spots fully matched	Spot fields correct		
100%	87.4%	59%	95.5%		
		63 of 107 reference spots	after pairing · 363 of 380		

SIGNALS FOR PRODUCT AND RESEARCH

Friction worth fixing

Each row is a reproducible issue, the partner impact, what we did, and the change that would remove it.

SIGNAL	WHAT BROKE / WHY IT MATTERS	WORKAROUND	ASK
P1 · PRODUCT Opaque Build errors	AI Studio Build returned permission across three models after a paid key was selected. Could not tell if the fault was key, project, model, or builder.	Built the app locally with the same Gemini key; kept AI Studio for prompt work.	Name the denied principal, missing permission, and the fix in the error.
P1 · API PDF quality knob missing	resolution / media_resolution on Interactions PDF items returned 400. Docs, SDK types, and service behavior disagreed.	Rendered pages at 180 DPI and sent images with resolution: high .	Align API, SDK, and docs. Ship one executable high-res PDF example.
P1 · EVAL Valid JSON, wrong value	Schema-valid output still chose the wrong business field and once auto-accepted a bad critical header. Shape checks alone hide failure modes.	Kept field checks, routing, label comparison, and audit outside the model.	Publish a reference that covers structured output, semantic scoring, and monitoring.
P1 · DATA Field contract vs labels	Invoice #, Order #, and Net Total were all plausible until precedence was written. Source audit found 5 conflicting headers and 64 corrupted spot labels.	Defined field precedence, recorded adjudications, excluded disputed labels from headline metrics.	Ship field definitions, label provenance, and quality flags with extraction benchmarks.

RESEARCH TAKEAWAY

Stop the example at “returns JSON.” Show how output is checked, routed, scored by failure mode, and monitored after accept.